

Logistic_Regression

December 30, 2023

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.simplefilter("ignore")
```

```
[3]: df=pd.read_csv("file:///Users/chaitanya_offical/Desktop/SRK-Dataset/diabetes.
↪csv")
```

```
[4]: df.head()
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null    int64
1   Glucose                768 non-null    int64
2   BloodPressure          768 non-null    int64
3   SkinThickness          768 non-null    int64
```

```

4   Insulin                768 non-null    int64
5   BMI                    768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                    768 non-null    int64
8   Outcome                768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB

```

```
[7]: df["Outcome"].unique()
```

```
[7]: array([1, 0])
```

```
[8]: df["Outcome"].value_counts()
```

```

[8]: 0    500
     1    268
     Name: Outcome, dtype: int64

```

```

[10]: continuous=["Glucose","BloodPressure","SkinThickness","Insulin","BMI","DiabetesPedigreeFunction"]
      discrete_count=["Pregnancies"]
      discrete_categorical=["Outcome"]

```

```
[11]: df[continuous].describe()
```

```

[11]:      Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
count  768.000000    768.000000    768.000000  768.000000  768.000000
mean   120.894531    69.105469    20.536458    79.799479    31.992578
std     31.972618    19.355807    15.952218   115.244002     7.884160
min      0.000000     0.000000     0.000000     0.000000     0.000000
25%     99.000000    62.000000     0.000000     0.000000    27.300000
50%    117.000000    72.000000    23.000000    30.500000    32.000000
75%    140.250000    80.000000    32.000000   127.250000    36.600000
max    199.000000   122.000000    99.000000   846.000000    67.100000


      DiabetesPedigreeFunction  Age
count              768.000000  768.000000
mean                   0.471876   33.240885
std                    0.331329   11.760232
min                    0.078000   21.000000
25%                   0.243750   24.000000
50%                   0.372500   29.000000
75%                   0.626250   41.000000
max                   2.420000   81.000000

```

```

[12]: sns.heatmap(df[continuous].corr(),annot=True)
      plt.show()

```



```
[13]: df.isnull().sum()
```

```
[13]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

```
[14]: X=df.drop('Outcome',axis=1)
      y=df['Outcome']
```

```
[15]: #Train/Test Split

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.
↳2,random_state=16)
```

```
[18]: from sklearn.preprocessing import StandardScaler

# Create an instance of StandardScaler
sc = StandardScaler()

# Apply StandardScaler to the continuous features in the training set
X_train.iloc[:, 1:8] = sc.fit_transform(X_train.iloc[:, 1:8])

# Apply the same StandardScaler to the continuous features in the test set
X_test.iloc[:, 1:8] = sc.transform(X_test.iloc[:, 1:8])
```

```
[20]: #Modelling

from sklearn.linear_model import LogisticRegression
log_reg=LogisticRegression()
log_reg.fit(X_train,y_train)
```

```
[20]: LogisticRegression()
```

```
[21]: # Assuming you have already trained a logistic regression model (log_reg) on
↳your training data

# Make predictions on the training set
ypred_train = log_reg.predict(X_train)

# Train accuracy
from sklearn.metrics import accuracy_score # Corrected the function name
print("Train Accuracy:", accuracy_score(y_train, ypred_train)) # Corrected the
↳function name

# Cross-validation score
from sklearn.model_selection import cross_val_score
print("CV Score:", cross_val_score(log_reg, X_train, y_train, cv=5,
↳scoring='accuracy')) # Specify the scoring metric
```

Train Accuracy: 0.7719869706840391

CV Score: [0.73170732 0.67479675 0.79674797 0.78861789 0.7704918]

```
[22]: ypred_test=log_reg.predict(X_test)
print("Test Accuarcy",accuracy_score(y_test,ypred_test))
```

Test Accuarcy 0.33766233766233766

```
[24]: #confusion matrix
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,ypred_test)
cm
```

```
[24]: array([[ 0, 102],
           [ 0,  52]])
```


```
[28]: from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Assuming you have the confusion matrix (cm) and display labels
display_labels = df["Outcome"].unique()

# Create a ConfusionMatrixDisplay object
cm_display = ConfusionMatrixDisplay(confusion_matrix=cm,
    ↪display_labels=display_labels)

# Plot the confusion matrix
cm_display.plot()

# Display the plot
plt.show()
```



```
[30]: from sklearn.metrics import classification_report
```


```
# Assuming y_test and ypred_test are your true and predicted labels  
print(classification_report(y_test, ypred_test))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	102
1	0.34	1.00	0.50	52
accuracy			0.34	154
macro avg	0.17	0.50	0.25	154
weighted avg	0.11	0.34	0.17	154

```
[31]: from sklearn.metrics import roc_curve, auc, RocCurveDisplay
```

```
fpr, tpr, thresholds = roc_curve(y_test, ypred_test)  
roc_auc = auc(fpr, tpr)  
RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=roc_auc).plot()
```

```
[31]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fb04b00ae20>
```



```
[32]: from sklearn.metrics import roc_auc_score  
      roc_auc_score(y_test,ypred_test)
```

```
[32]: 0.5
```

```
[ ]:
```